

Azure NoSQL Offerings


Eshant Garg

Azure Data Engineer, Architect, Advisor

eshant.garg@gmail.com

Why NoSQL DB?

What traditional databases were lacking?



An Introduction to
NoSQL

RDBMS were lacking




Scalability




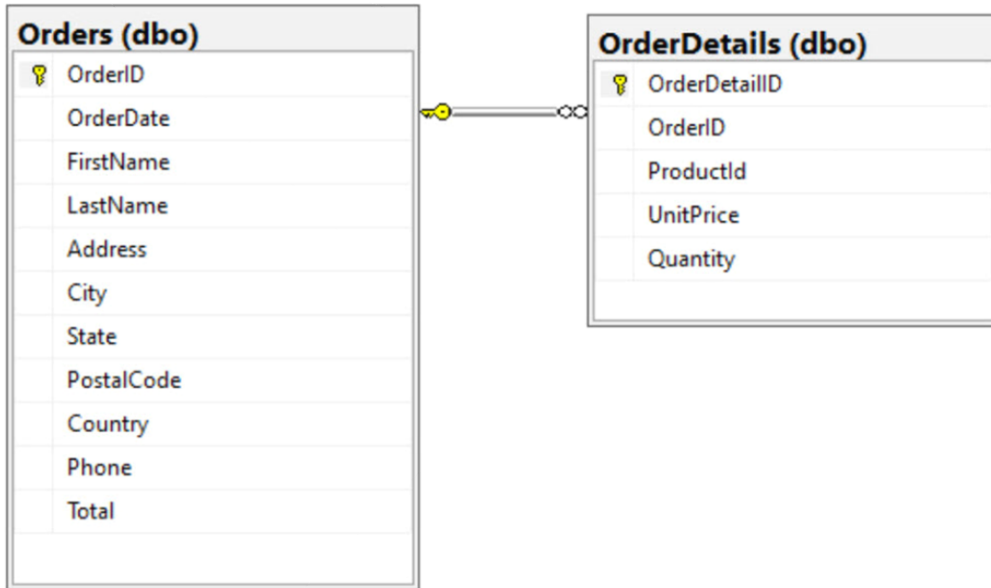
Flexibility

What is NoSQL



- **Vertical scaling**
 - Add more CPU, RAM, HDD in same system
- **Horizontal Scaling**
 - Add more commodity machines in system






```
{
  "OrderId": 1,
  "OrderDate": 1574161910220,
  "FirstName": "John",
  "LastName": "Smith",
  "Address": "10 Street",
  "City": "City",
  "State": "VA",
  "OrderDetails": [
    {
      "UnitPrice": 7.99,
      "OrderDetailId": 2,
      "Quantity": 1,
      "ProductId": 259694,
      "OrderId": 1
    },
    {
      "UnitPrice": 7.99,
      "OrderDetailId": 3,
      "Quantity": 1,
      "ProductId": 295693,
      "OrderId": 1
    }
  ],
  "id": "795c50dc-1a83-11ea-bf07-00163ee85f66",
  "_rid": "VdgtAK23OMANAAAAAAAAA==",
  "_self": "dbs/VdgtAA==/colls/VdgtAK23OMA=/docs/VdgtAK23OMANAAAAAAAAA==/",
  "_etag": "\"370017e1-0000-1100-0000-5df770f20000\"",
  "_attachments": "attachments/",
  "_ts": 1576497394
}
```

```
{
  "orderid": 12212,
  "orderdate": "12/4/2020",
  "customer":
    { "name": "Bob Smith", "email": "bobsmith@email.bob" },
  "status": "in process",
  "paymentmethod": "invoice",
  "products": [
    { "name": "Product 1", "quantity": 1 },
    { "name": "Product 2", "quantity": 1, status: 3 }
  ]
}
```

NoSQL Use Cases



NoSQL Limitations

No SQL

Schema-less data means inconsistent data

Denormalized data means redundant data

Redundant data means inaccuracies and conflicts

Does not support many good features of Relational DB

- SPs, Functions, Views, Row level security, Locks, etc.

SQL vs NoSQL

SQL


- Relational database
- Fixed schema
- Designed for complex queries
- SQL, MySql, Oracle, Postgres
- Vertical scaling
- Row Oriented
- Tables
- Limited for big data

NoSQL


- Non-relational or distributed
- Dynamic
- Not for complex queries
- MongoDB, Redis, Hbase
- Horizontal scaling
- Multi-model oriented
- Collections
- Great for big data

4 Types of NoSQL Databases


Document




Graph



Key-Value




Wide-column



Key-value store

Key-Value




Phone Directory

Key	Value
Bob	(123) 456-7890
Jane	(234) 567-8901
Tara	(345) 678-9012
Tiara	(456) 789-0123

- Uses a simple key/value to store data
- Quick to query due to its simplicity
- Value can be JSON, BLOB, String etc.
- **Use Cases:**
 - User profiles and session info on a website, blog comments, telecom directories, IP forwarding tables, shopping cart contents on e-commerce sites, and more.
- **Examples**
 - Cosmos DB Table API, Redis, Table Storage, Oracle NoSQL Database, Voldemorte, Aerospike, Oracle Berkeley DB

Document store

Document



- Document-oriented model to store data
- Similar to key/value store, difference is that, the value in a document store database consists of semi-structured data.
- Each record and its associated data within a single document.
- Document stores are usually XML, JSON, BSON, YAML, etc.
- **Use Cases:**
 - Content management systems, blogging platforms, and other web applications, blog comments, chat sessions, tweets, ratings, etc.

- **Examples**

- Cosmos DB, MongoDB, DocumentDB, CouchDB, MarkLogic, OrientDB

```
{
  "orderid": 12212,
  "orderdate": "12/4/2020",
  "customer":
    { "name": "Bob Smith", "email": "bobsmith@email.bob" },
  "status": "in process",
  "paymentmethod": "invoice",
  "products": [
    { "name": "Product 1", "quantity": 1 },
    { "name": "Product 2", "quantity": 1, status: 3 }
  ]
}
```

Column store



UserProfile

Bob	emailAddress	gender	age
	bob@example.com	male	35
	1465676582	1465676582	1465676582
Britney	emailAddress	gender	
	brit@example.com	female	
	1465676432	1465676432	
Tori	emailAddress	country	hairColor
	tori@example.com	Sweden	Blue
	1435636158	1435636158	1465633654

- Stores data using a column oriented model
- Columns in each row are contained within that row
- Each row can have different columns to the other rows.
- Extremely quick to load and query
- **Use Cases:**
 - Sensor Logs [Internet of Things (IOT)], User preferences, Geographic information, Reporting systems, Time Series Data, Logging and other write heavy applications
- **Examples**
 - Cosmos DB, Bigtable, Cassandra, Hbase, Vertica, Druid, Accumulo, Hypertable

Graph store

Graph




- Focuses on how data relates to other data points.
- A **node** is a specific entity or piece of information
- **Edge** simply specifies the relationship between two nodes.
- **Use Cases:**
 - Social networks, realtime product recommendations, network diagrams, fraud detection, access management, and more.
- **Examples**
 - Cosmos DB Gremlin API, Neo4j, Blazegraph, and OrientDB.

Multi-model

- Include features/characteristics of more than one data model.
- **Example:**
 - **OrientDB:** OrientDB combines a graph model with a document model.
 - **ArangoDB:** Uses key/value, document, and graph models.
 - **Virtuoso:** Combines relational, graph, and document models.

NoSQL Offerings by Microsoft Azure



Microsoft Azure Storage

Highly Available


Durable

Secure

Scalable

Cost effective

Accessible



Programmatic Access to Storage Accounts

REST APIs

SDKs

PowerShell

Azure CLI

Azure Storage
Explorer

AzCopy

Azure Storage Account Type

General Purpose V2

Supported Services: Blob, File, Disk, Table and Queue Storage

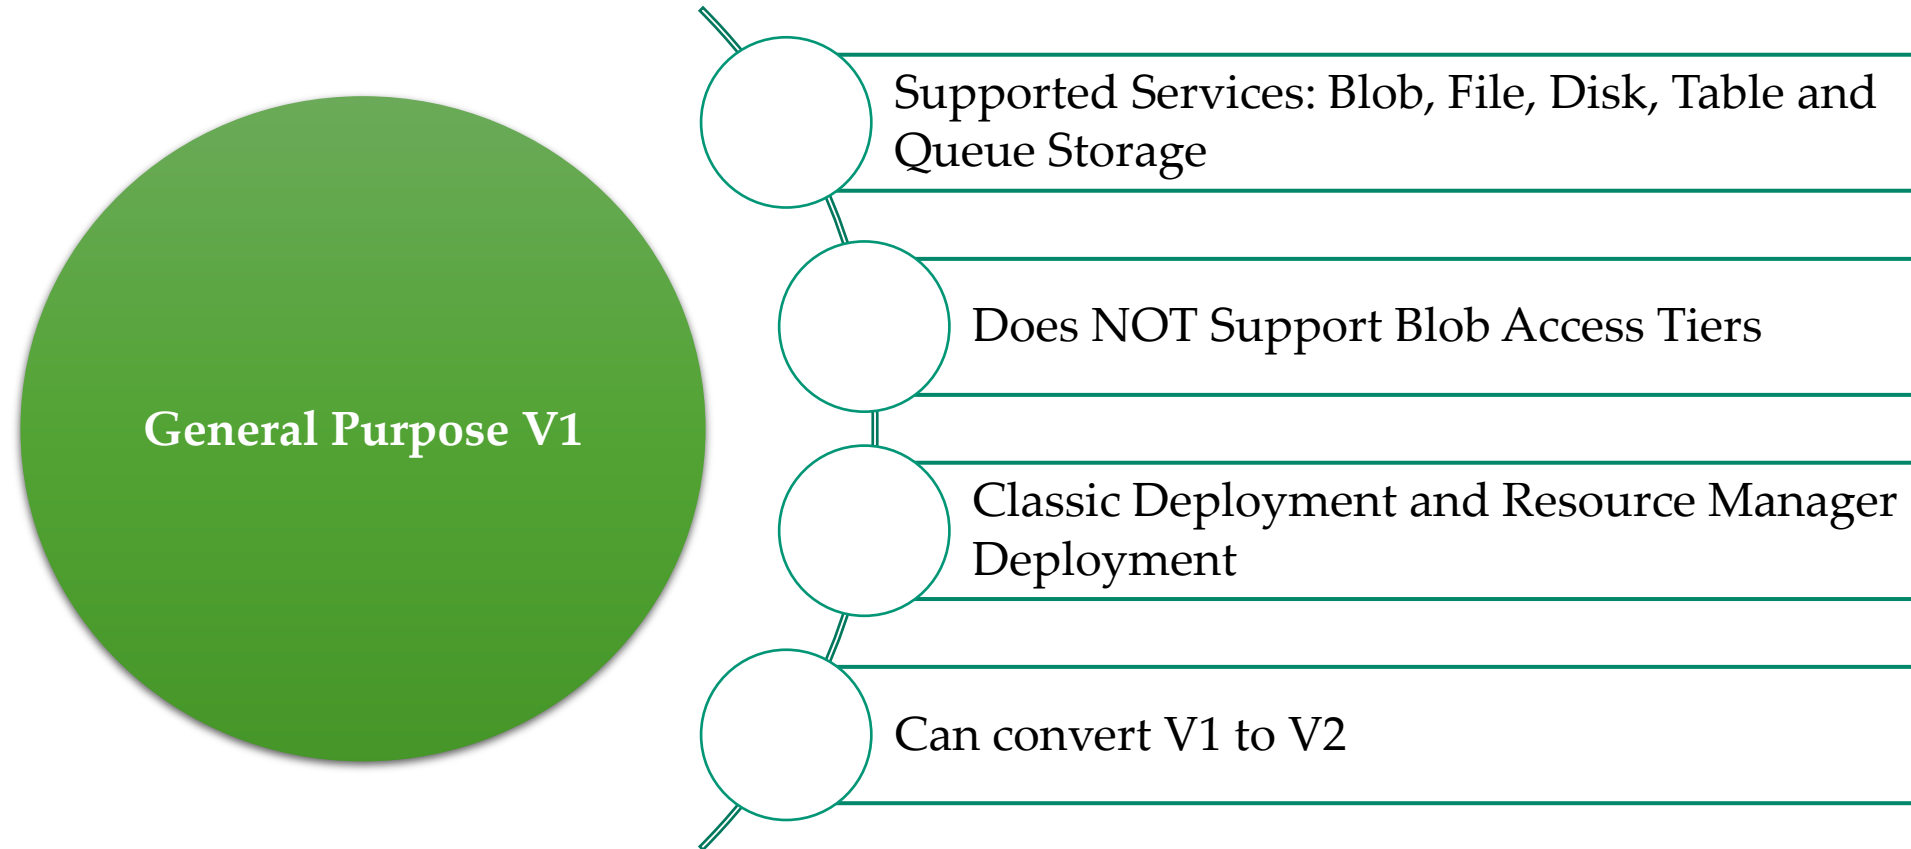
Supports Blob Access Tiers

Block Blobs, Append Blobs, Page Blobs


Hierarchical namespace support (Data Lake Gen2)

Premium tier available for Page Blobs only

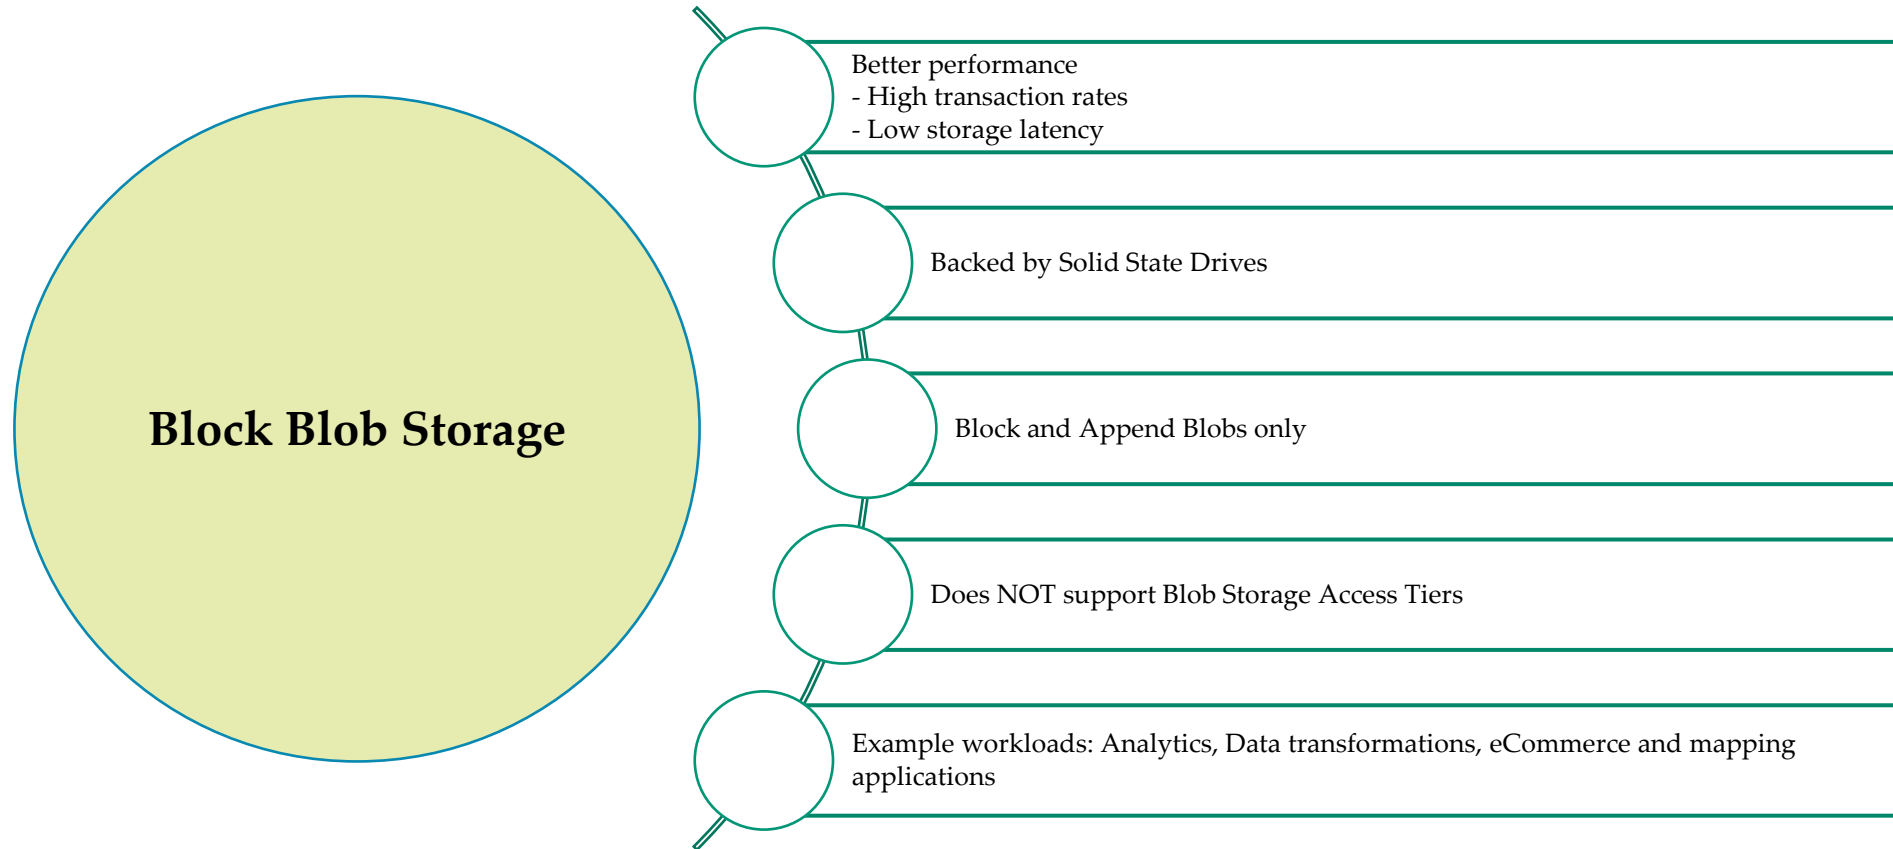
Azure Storage Account Type



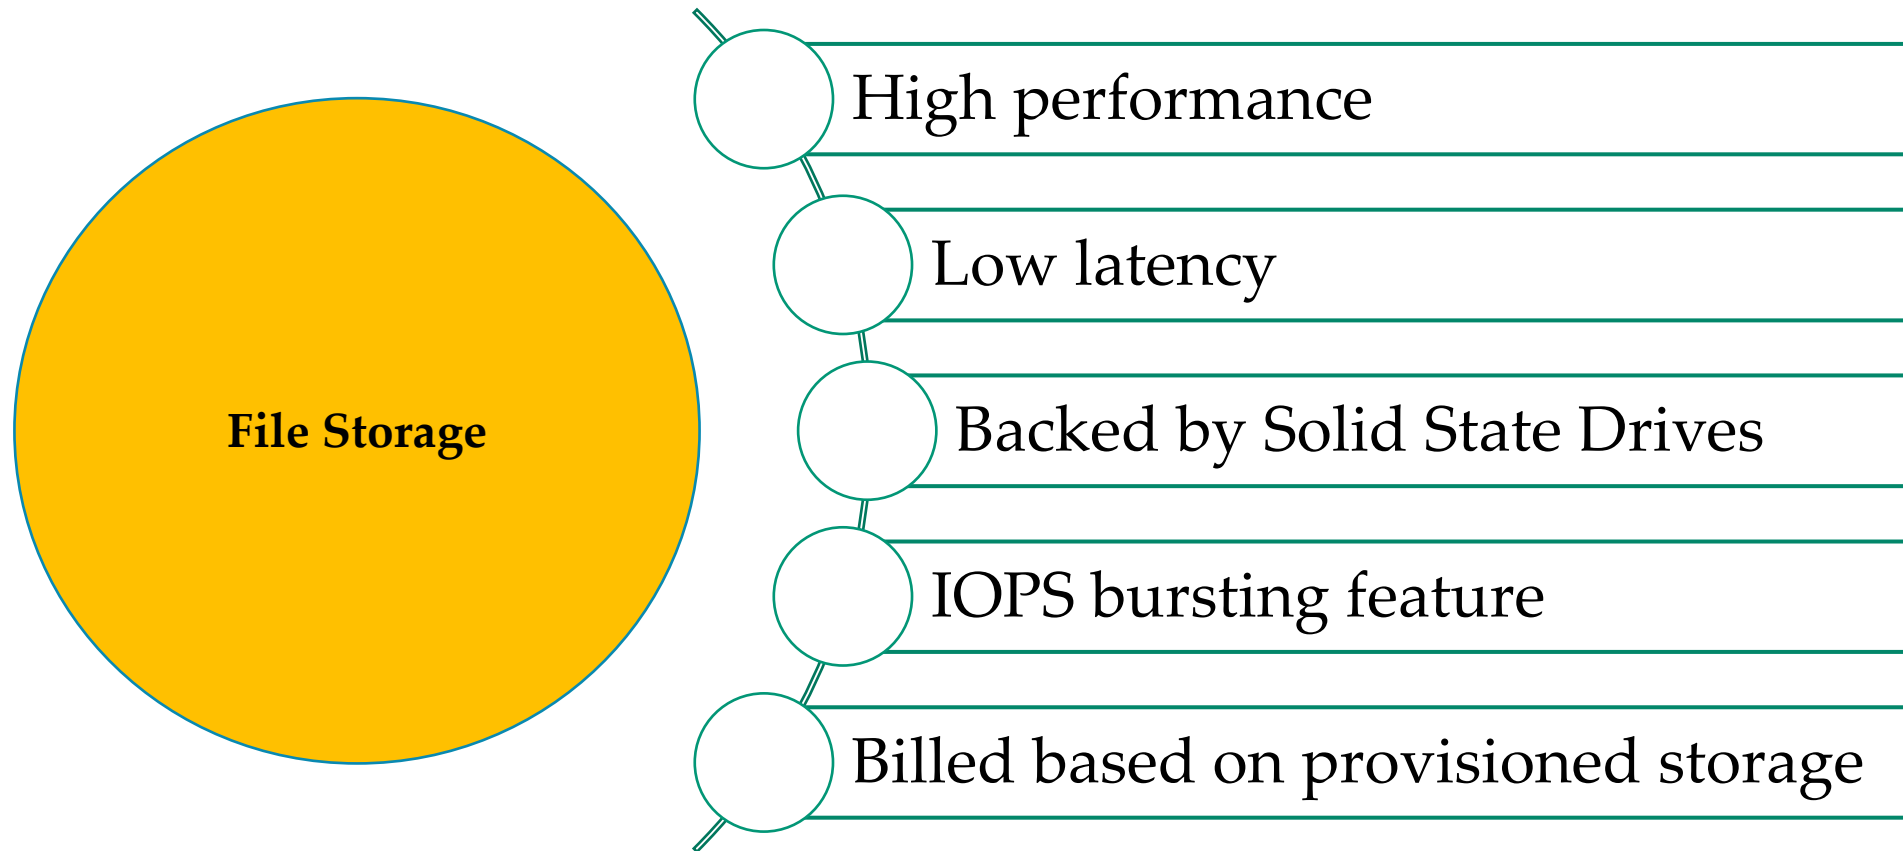
Azure Storage Account Type



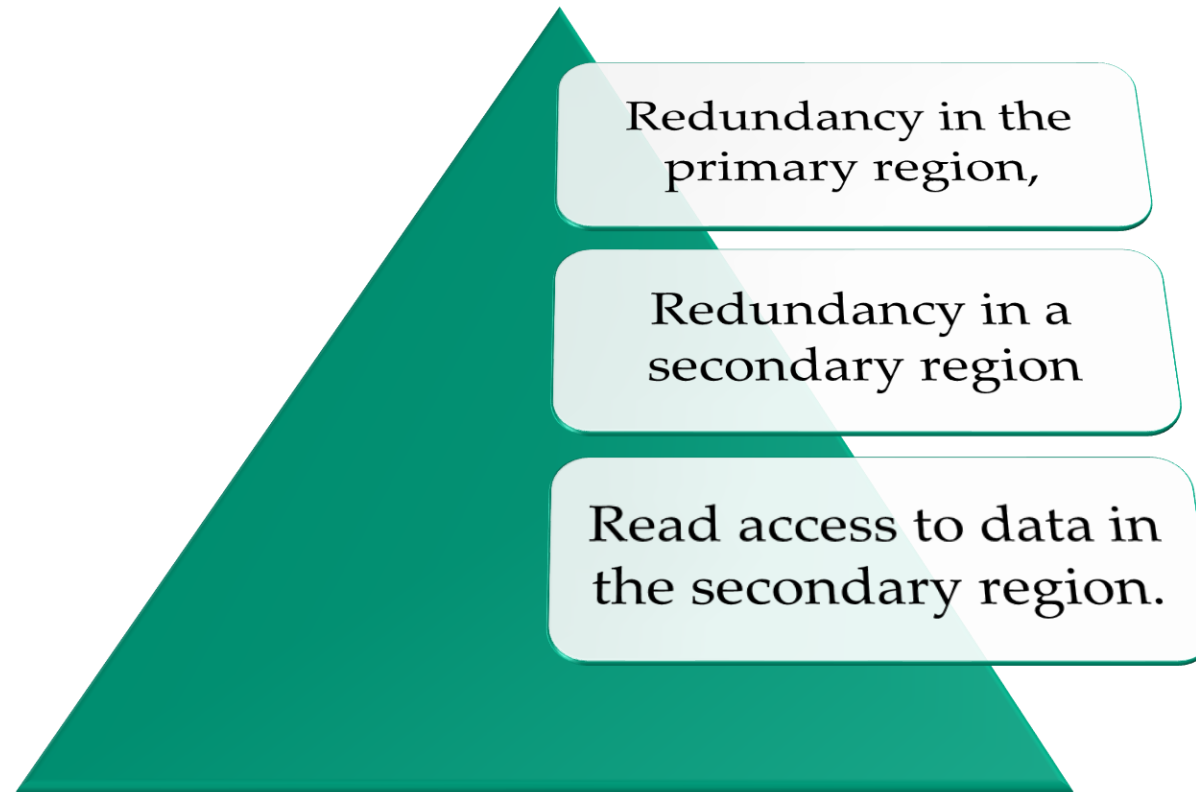
Azure Storage Account Type



Azure Storage Account Type





Three categories of replication options



Locally Redundant Storage (LRS)


Region A



Storage Clusters

Each cluster is physically separate in what's called an availability zone, with its own separate utilities and networking.

Region B




Region B

Hundreds of miles away from the primary region to prevent data loss in the event of a natural disaster.

Zone Redundant Storage (ZRS)

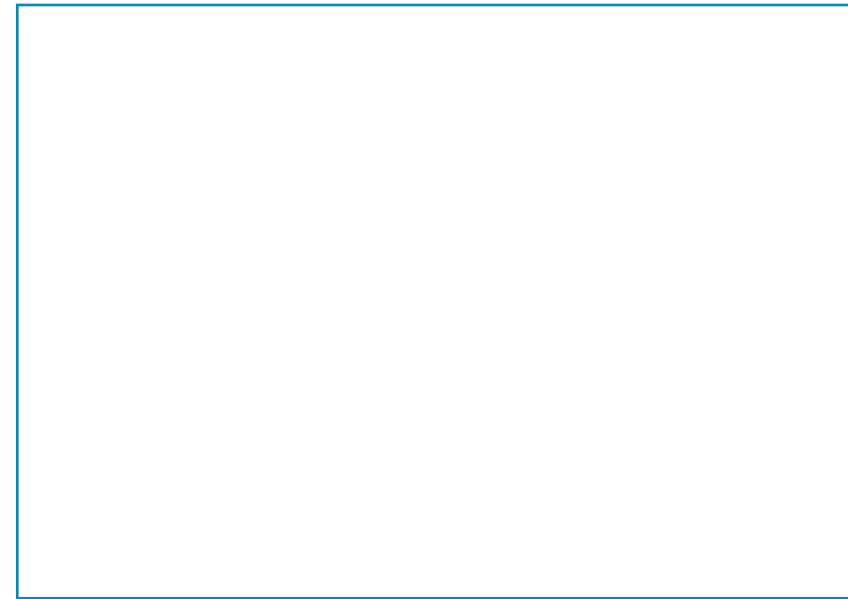
Region A



Storage Clusters

Each cluster is physically separate in what's called an availability zone, with its own separate utilities and networking.

Region B




Region B

Hundreds of miles away from the primary region to prevent data loss in the event of a natural disaster.

Geo Redundant Storage (GRS)


Region A



Storage Clusters

Each cluster is physically separate in what's called an availability zone, with its own separate utilities and networking.

Region B




Region B

Hundreds of miles away from the primary region to prevent data loss in the event of a natural disaster.

Geo Zone Redundant Storage (GZRS)


Region A



Storage Clusters

Each cluster is physically separate in what's called an availability zone, with its own separate utilities and networking.

Region B




Region B

Hundreds of miles away from the primary region to prevent data loss in the event of a natural disaster.

Read access geo Redundant Storage (RA-GRS)


Region A



Storage Clusters

Each cluster is physically separate in what's called an availability zone, with its own separate utilities and networking.

Region B (Read)




Region B

Hundreds of miles away from the primary region to prevent data loss in the event of a natural disaster.

Read access Geo Zone Redundant Storage (RA-GZRS)


Region A



Storage Clusters

Each cluster is physically separate in what's called an availability zone, with its own separate utilities and networking.


Region B (Read)




Region B

Hundreds of miles away from the primary region to prevent data loss in the event of a natural disaster.

Blob Access Tiers



Highest storage cost
Lowest data access cost



Lowest storage cost
Higher data access cost



Lowest storage cost
Highest data retrieval cost
Data is offline

Azure Blob Storage Lifecycle Management

A large blue arrow pointing to the right, spanning the width of the diagram.

Azure Blob Storage

- Designed for images and unstructured Data
 - Store Documents and access in browser
 - Database backup
 - Store audio and video files and stream them
 - Store data for analysis
 - Log files
- Scalability
- Cheapest way to store data in azure
- Simple design and easy to use
- HDFS and blob storage REST APIs

Microsoft Azure
Blob Storage



Blob Types

- **Block Blob**
 - Composed for Blocks
- **Append Blob**
 - Can only append blocks
 - Ideal for logs
- **Page Blob**
 - VM disks and databases
 - Frequent random read/write applications

Microsoft Azure
Blob Storage



Use cases

- Only basic storage is needed
- Data is unstructured
- Data that is older or not used as much
- Money is an issue

Microsoft Azure
Blob Storage



Advantages of Blob storage

- Extremely cheap
- Simple to setup
- No configuration
- Doesn't require powerful computing to manage

Microsoft Azure
Blob Storage



Limitations of Blob storage

- No Indexes
- No Search Tools
- Not optimized for performance
- You are responsible for replication and synchronization
- Requires external compute to process

Microsoft Azure
Blob Storage

